



Beag Labs

\$4.5M PRE-SEED ROUND

Helping enterprises own their models and inference.

Small models, deployed anywhere. Beag Labs builds domain-specific AI for internal enterprise workflows — CRM, compliance, IAM, tool calling — that runs on your infrastructure at a fraction of the cost.

MEET THE FOUNDER

Raising \$4.5M pre-seed.

\$200B

Enterprise AI market by 2027. The majority of spend will shift from general-purpose APIs to domain-specific, deploy-anywhere models.

13x cheaper

Small models outperform GPT-4 on internal enterprise workflows at 13x lower inference cost and 40x lower latency. CRM, compliance, IAM, tool calling — measured on real customer workloads.

THE MARKET PROBLEM

Enterprise AI is stuck between overpriced APIs and impossible build-vs-buy decisions.

Every company needs AI for internal enterprise workflows. No one sells a model that is accurate enough, private enough, and cheap enough — so enterprises either overpay for general APIs or give up.

01

Internal workflows are stuck on general-purpose APIs

Companies pay per-token rates for 1.8-trillion-parameter models to run CRM updates, compliance checks, IAM provisioning, and workflow automation. These internal tasks need a focused model that costs 13x less and runs 40x faster — but no one sells one.

02

Regulated industries can't use the cloud

Legal, healthcare, defense, and financial services cannot send sensitive data to third-party inference APIs. Every compliance officer's first question is 'where does this data go?' — and every vendor's answer is unacceptable.

03

Custom models are too expensive to build

Fine-tuning requires ML engineers, labeling budgets, and GPU infrastructure that most teams don't have. The result: enterprises either overpay for general APIs or give up on AI entirely.

HOW WE WIN

Enterprises want to own their models and inference. We own the pipeline that builds them.

We sell the outcome, not the API call. Our pipeline uses frontier base SLMs (Google Gemma E4B) as the base for domain-specific fine-tuning of internal models deployed within the customer's infrastructure.

01

Open weights, fine-tuned for any workflow

We start with frontier open-weight models (Google Gemma E4B) and fine-tune them for the customer's internal workflows. The resulting model is deployed on their infrastructure — on-prem, VPC, or edge device. The customer owns the model weights under a standard commercial license.

02

GRPO + OPD = Small Models that run with the big dogs

Our training recipe uses GRPO and On-policy Distillation to produce small models that can outperform foundation models on internal workflows. We measure performance on held-out test sets and compare against GPT-4 and Claude. The result: 13x lower inference cost, 40x lower latency, and better accuracy on the customer's internal tasks.

03

Customer data never touches our servers after training

We train on your data, in your cloud account or air-gapped environment if needed. The resulting model deploys on your infrastructure — on-prem, VPC, or edge device. We never see your inference data. You own the model weights under a standard commercial license.

04

Each deployment compounds into the next

Every engagement teaches us which schemas, evaluation patterns, and data transforms generalize. Our platform gets more productive with each customer instead of starting from zero. The models stay with the customer; the platform knowledge stays with us.

TRACTION & BUSINESS MODEL

Pre-revenue with measurable technical traction and pipeline.

We have not yet recognized revenue — we are pre-revenue by design. What we do have is a validated technical approach, active design partners in legal and healthcare, and model performance numbers that beat GPT-4 and Claude on their specific enterprise workflow tasks.

REVENUE

Pre-revenue

Founder-led delivery engagements in progress. First paid contracts expected Q3 2026.

DESIGN PARTNERS

3 active

Working with legal-tech and healthcare organizations on production workflow automation pipelines. Each is a paid pilot with a path to multi-year contract.

MODEL PERFORMANCE

92-96% F1

Domain-specific workflow models consistently outperform GPT-4 and Claude on held-out test sets at 13x lower inference cost.

TARGET CONTRACT

\$150-250k ACV

2-3 year engagements for workflow-level outcomes, not model-as-a-service token meters. Customers buy a result, not an API key.

REVENUE MODEL

Phase 1: Outcomes-based delivery

Multi-year contracts priced on workflow scope, not token volume. Customer owns the model. We own the pipeline. \$150-250k ACV.

Phase 2: Platform self-serve

As the pipeline productizes, customers can run their own workflow configurations and deployments. Platform subscription + usage-based deployment pricing.

Phase 3: Vertical expansion

Pre-built models for CRM, compliance, IAM, and workflow automation. Each new customer in a vertical benefits from every previous customer's data and schemas.

OPERATING SPLIT

Delivery & workflow ownership 70%



Funds operations and grows reference customer base

Platform R&D 30%



Turns delivery patterns into product features

Services revenue funds the R&D that makes services cheaper and faster. Each engagement fully covers its own cost plus a margin that reinvests into the platform. Neither track starves the other.

DEFENSIBILITY

Four compounding advantages that get stronger with every deployment.

This is not a features race. Our moat is structural — built into the delivery model and data flywheel.

MOAT

- 01 The uncertainty engine gets smarter with every deployment. The distribution of frontier-model confidence across enterprise workflows is proprietary data that improves our pipeline efficiency permanently.
- 02 Customer-owned deployment creates structural switching costs. Your legal team reviewed the outputs, your IT deployed the model, and your compliance team approved the air-gap. Replacing the model means redoing all three.
- 03 Vertical knowledge compounds. CRM schemas, compliance taxonomies, IAM role structures, workflow orchestration patterns — these don't exist in public training data. Each engagement creates defensible IP for that vertical.
- 04 The 70/30 split keeps us grounded. Delivery revenue funds research; research makes delivery cheaper and faster. Neither track starves the other. This is a durable business model, not a burn-for-market-share play.

WHO WE COMPETE WITH

General-purpose APIs (GPT-4, Claude, Gemini)

Too expensive for volume workflow automation. Cannot be deployed on customer infrastructure. No data privacy guarantees.

Fine-tuning platforms (Tinker, Fireworks, Together)

Sell GPU access, not workflow outcomes. Customer still needs ML engineers, labeling pipelines, and evaluation infrastructure. We absorb that complexity.

In-house ML teams

Most enterprises cannot hire and retain the talent needed to build and maintain custom models. We act as their ML department with a productized delivery model.

ROADMAP

Each phase de-risks the next. Revenue funds the roadmap.

We are capital-efficient by necessity and design. Revenue from delivery funds platform R&D. The platform R&D makes delivery cheaper and faster. This is a flywheel, not a burn multiple.

PHASE 1

Now

STEP 01

Land and deliver — prove the model

Win 4-6 high-trust design partners through founder-led sales. Each engagement builds tooling, evaluation infrastructure, and a reference customer. Revenue funds the roadmap. Target: \$600k ARR exiting year one.

PHASE 2

Next

STEP 02

Productize the pipeline

Turn repeated delivery patterns into self-serve tooling: data connectors, the uncertainty-driven labeling interface, one-click deployment to customer infrastructure. Each workflow deployment gets faster and cheaper to fulfill.

PHASE 3

Scale

STEP 03


Expand vertically

Each vertical (legal, healthcare, finance, defense) has a compounding data and schema advantage. Models, evaluation datasets, and workflow templates transfer across customers within a vertical. Enter new verticals through strategic design partners.

FINANCIAL PLAN

\$4.5M pre-seed. Training infrastructure is the largest line item — everything else stays lean.

Capital efficiency is the thesis. We allocate the majority of funds to model training infrastructure because that is the core product. Everything else — team, operations, legal — is minimized by design.



WHY TRAINING IS #1

Core product = trained models
Every dollar in training infrastructure produces a deployed model that generates revenue. This is not overhead — it is the product.

Iteration speed = competitive advantage
8x B200 compute means full training cycles in 3-5 hours. Faster iteration = better models = faster revenue.


Cost per training run
~\$200-300 per complete training cycle (SFT + GRPO + OPD). Each run produces a production-ready model. Budget supports 5,000+ training iterations.

REVENUE FORECAST

Capital-efficient path to \$2.5M ARR within 18 months of first contract.

Forecast assumes founder-led delivery through year one with GTM and Sales Engineer hires triggered by pipeline milestones. Revenue begins when the first design partners convert to paid contracts in Q3. Cash runway extends past 18 months even at conservative revenue projections.

MONTHLY REVENUE VS. EXPENSES (\$M)



REVENUE DRIVERS

4-6 design partners convert to paid contracts in year one at \$150-250K ACV. Revenue compounds as each deployment validates the model for adjacent prospects.

EXPENSE DISCIPLINE

Burn rate stays under \$280K/mo at peak. Founder operates at 5x capacity via LLMs. Hires are gated on revenue, not planning assumptions.

BREAKEVEN TRAJECTORY

Revenue catches expenses around month 22. At target case, the company reaches operational breakeven before the full cash runway is consumed.

TEAM STRATEGY

Solo founder. LLM-powered operations. Hires only where revenue demands it.

We use LLMs to replace functions that would otherwise require 3-5 engineering hires. The founder handles product, ML, engineering, and delivery. Every hire must directly generate or close revenue.

WHAT LLMs HANDLE TODAY

- ✓ **Code review & refactoring** — Claude / Cursor for iterative development
- ✓ **Documentation & content** — Generated, not hand-written
- ✓ **Testing & eval pipelines** — LLM-written test suites, automated evals
- ✓ **UI/UX design iterations** — AI-assisted prototyping and iteration
- ✓ **Customer support triage** — Automated responses, founder handles escalation
- ✓ **Data labeling & annotation** — Frontier models label 95% of training examples

Net effect: founder operates at the capacity of a 5-person team without the coordination overhead or burn rate.

PLANNED HIRES (REVENUE-GATED)

GTM / Revenue Lead

Month 6-9

Owns pipeline, manages inbound/outbound, runs discovery calls. Founder stays on demos and technical close. This hire is triggered by pipeline volume — not a planning assumption.

Sales Engineer

Month 12-15

Runs POC deployments, handles technical due diligence, builds custom demos. Hired only when deal volume exceeds founder capacity. Directly tied to closing revenue.

Key principle: No engineering hires in year one. LLMs handle code, tests, docs, and design. First engineering hire (if needed) is year two, funded entirely by revenue.


Model sizes

Two model tiers designed for different deployment realities. One for edge and VPC. One for large enterprise operations.

4B PARAMS

Beag Labs Starling Satellite

Edge & VPC deployment. Fine-tuned Gemma E4B for CRM automation, compliance workflows, IAM provisioning, and tool-call orchestration.



- CRM, compliance, IAM, tool calling
- Deploys on-prem, VPC, or edge
- SFT + GRPO + On-Policy Distillation
- ~\$200-300 per training run

LARGE MODEL

Beag Labs Starling Mothership

Large enterprise operations model. Multi-step workflow orchestration across CRM, compliance, IAM, and cross-system tool calling.



- Multi-system workflow orchestration
- Cross-domain CRM, compliance, IAM
- Complex tool-call chains
- Training plan to be defined

THE ASK

\$4.5M pre-seed. 18 months of runway. Build the platform while customers pay for delivery.

We are looking for aligned capital that understands services-led SaaS and believes the enterprise AI market will reward focused, deploy-anywhere models over bloated general-purpose APIs. We will not grow at all costs — we will grow at the speed our customers pay us to.

USE OF FUNDS

- 01 Training infrastructure:** Baseten SLURM compute, 8x B200 training runs, data pipelines, evaluation infrastructure. This is the core product — every dollar here produces a deployed model that generates revenue.
- 02 Operations budget:** Founder salary, 2 revenue-gated hires (GTM/Revenue Lead month 6-9, Sales Engineer month 12-15), legal, compliance, insurance. No engineering hires in year one — LLMs handle code, tests, docs, and design.
- 03 Cloud & deployment:** Customer-funded cloud spend offset, staging environments, CI/CD. Cloud costs are structurally offset by customer contracts.

IDEAL PARTNER

- 01** Believes services-led go-to-market is a viable wedge into enterprise SaaS, not a consulting trap.
- 02** Understands regulated industries (legal, healthcare, defense) and why they cannot use cloud AI APIs.
- 03** Prefers durable company formation with honest metrics over growth-stage revenue fabrication.
- 04** Has the network to open design partner conversations in legal, healthcare, or financial services.

SCHEDULE A CONVERSATION